



| |
|---|
| Experiment No. 3 |
| Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model |
| Date of Performance: |
| Date of Submission: |

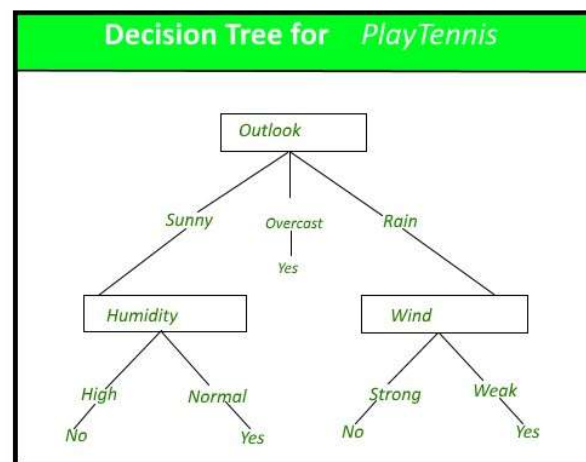


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked. fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool. education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Marriedspouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspect, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black. sex:

Female, Male.

capital-gain: continuous. capital-loss:

continuous. hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, OutlyingUS(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines,

Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.



Conclusion:

1. Dealing with Categorical Attributes during Data Pre-processing:

In the code, categorical attributes are handled as follows:

Label Encoding: Categorical attributes are transformed into numerical format using label encoding. The LabelEncoder from Scikit-Learn is applied to convert categorical values into unique integers.

Dropping Missing Values: Rows containing missing values in categorical columns ('workclass', 'occupation', 'native.country') are dropped from the dataset to maintain data quality.

2. Hyper parameter Tuning:

Hyperparameter tuning is minimal in the code. The Decision Tree classifier is created with a max_depth of 5. However, there's no comprehensive hyperparameter tuning process, such as grid search or random search, to optimize the model's performance. Tuning hyperparameters like max_depth, min_samples_split, and min_samples_leaf could potentially improve the model's accuracy.

3. Accuracy, Confusion Matrix, Precision, Recall, and F1-Score:

Accuracy: Measures overall correctness of predictions.

Confusion Matrix: Summarizes true positives, true negatives, false positives, and false negatives.

Precision: Measures accurate positive classifications.

Recall (Sensitivity): Measures identifying relevant instances.

F1-Score: Balances precision and recall, especially useful for imbalanced classes.

```
confusion matrix
[[ 764  735]
 [ 214 4315]]
precision    recall  f1-score   support

0           0.78     0.51     0.62     1499
1           0.85     0.95     0.90     4529

accuracy          0.84     6028
macro avg         0.82     0.73     0.76     6028
weighted avg      0.84     0.84     0.83     6028
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv("adult.csv")
df.head()
```

| | age | workclass | fnlwgt | education | education.num | marital.status | occupation | relatio |
|---|-----|-----------|--------|--------------|---------------|----------------|-------------------|---------|
| 0 | 90 | ? | 77053 | HS-grad | 9 | Widowed | ? | Not-in- |
| 1 | 82 | Private | 132870 | HS-grad | 9 | Widowed | Exec-managerial | Not-in- |
| 2 | 66 | ? | 186061 | Some-college | 10 | Widowed | ? | Unnr |
| 3 | 54 | Private | 140359 | 7th-8th | 4 | Divorced | Machine-op-inspct | Unnr |
| 4 | 41 | Private | 264663 | Some-college | 10 | Separated | Prof-specialty | Ow |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country        32561 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df.describe()
```

```
df.describe()
```

| | age | fnlwtg | education.num | capital.gain | capital.loss | hours.per.w |
|--------------|--------------|--------------|---------------|--------------|--------------|--------------|
| count | 32561.000000 | 3.256100e+04 | 32561.000000 | 32561.000000 | 32561.000000 | 32561.000000 |
| mean | 38.581647 | 1.897784e+05 | 10.080679 | 1077.648844 | 87.303830 | 40.437144 |
| std | 13.640433 | 1.055500e+05 | 2.572720 | 7385.292085 | 402.960219 | 12.345678 |
| min | 17.000000 | 1.228500e+04 | 1.000000 | 0.000000 | 0.000000 | 1.000000 |
| 25% | 28.000000 | 1.178270e+05 | 9.000000 | 0.000000 | 0.000000 | 40.000000 |
| 50% | 37.000000 | 1.783560e+05 | 10.000000 | 0.000000 | 0.000000 | 40.000000 |
| 75% | 48.000000 | 2.370510e+05 | 12.000000 | 0.000000 | 0.000000 | 45.000000 |
| max | 90.000000 | 1.484705e+06 | 16.000000 | 99999.000000 | 4356.000000 | 99.000000 |

```
# Info for categorical features
```

```
df.describe(include=['O'])
```

| | workclass | education | marital.status | occupation | relationship | race | sex | na |
|---------------|-----------|-----------|--------------------|----------------|--------------|-------|-------|-------|
| count | 32561 | 32561 | 32561 | 32561 | 32561 | 32561 | 32561 | 32561 |
| unique | 9 | 16 | 7 | 15 | 6 | 5 | 2 | 1 |
| top | Private | HS-grad | Married-civ-spouse | Prof-specialty | Husband | White | Male | |
| freq | 22696 | 10501 | 14976 | 4140 | 13193 | 27816 | 21790 | |

```
duplicated_rows = df.duplicated()
```

```
any_duplicates = duplicated_rows.any()
```

```
print("Duplicated Rows:")
```

```
df[duplicated_rows]
```

Duplicated Rows:

| | age | workclass | fnlwgt | education | education.num | marital.status | occupation | rel |
|--------------|-----|------------------|--------|--------------|---------------|--------------------|-------------------|-----|
| 8453 | 25 | Private | 308144 | Bachelors | 13 | Never-married | Craft-repair | N |
| 8645 | 90 | Private | 52386 | Some-college | 10 | Never-married | Other-service | N |
| 12202 | 21 | Private | 250051 | Some-college | 10 | Never-married | Prof-specialty | |
| 14346 | 20 | Private | 107658 | Some-college | 10 | Never-married | Tech-support | N |
| 15603 | 25 | Private | 195994 | 1st-4th | 2 | Never-married | Priv-house-serv | N |
| 17344 | 21 | Private | 243368 | Preschool | 1 | Never-married | Farming-fishing | N |
| 19067 | 46 | Private | 173243 | HS-grad | 9 | Married-civ-spouse | Craft-repair | |
| 20388 | 30 | Private | 144593 | HS-grad | 9 | Never-married | Other-service | N |
| 20507 | 19 | Private | 97261 | HS-grad | 9 | Never-married | Farming-fishing | N |
| 22783 | 19 | Private | 138153 | Some-college | 10 | Never-married | Adm-clerical | |
| 22934 | 19 | Private | 146679 | Some-college | 10 | Never-married | Exec-managerial | |
| 23276 | 49 | Private | 31267 | 7th-8th | 4 | Married-civ-spouse | Craft-repair | |
| 23660 | 25 | Private | 195994 | 1st-4th | 2 | Never-married | Priv-house-serv | N |
| 23720 | 44 | Private | 367749 | Bachelors | 13 | Never-married | Prof-specialty | N |
| 23827 | 49 | Self-emp-not-inc | 43479 | Some-college | 10 | Married-civ-spouse | Craft-repair | |
| 26738 | 23 | Private | 240137 | 5th-6th | 3 | Never-married | Handlers-cleaners | N |
| 27133 | 28 | Private | 274679 | Masters | 14 | Never-married | Prof-specialty | N |
| 28796 | 27 | Private | 255582 | HS-grad | 9 | Never-married | Machine-op-inspct | N |
| 29051 | 42 | Private | 204235 | Some-college | 10 | Married-civ-spouse | Prof-specialty | |

| | | | | | | | |
|-------|----|---------|-------|---------|---|--------------------|--------------|
| 29334 | 39 | Private | 30916 | HS-grad | 9 | Married-civ-spouse | Craft-repair |
|-------|----|---------|-------|---------|---|--------------------|--------------|

```
df = df.drop_duplicates()
```

```
# Correction of target value using a map
income_map = {'<=50K': 1, '>50K': 0}
df['income'] = df['income'].map(income_map)
```

<ipython-input-47-00c8c2884cd1>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/user>
df['income'] = df['income'].map(income_map)

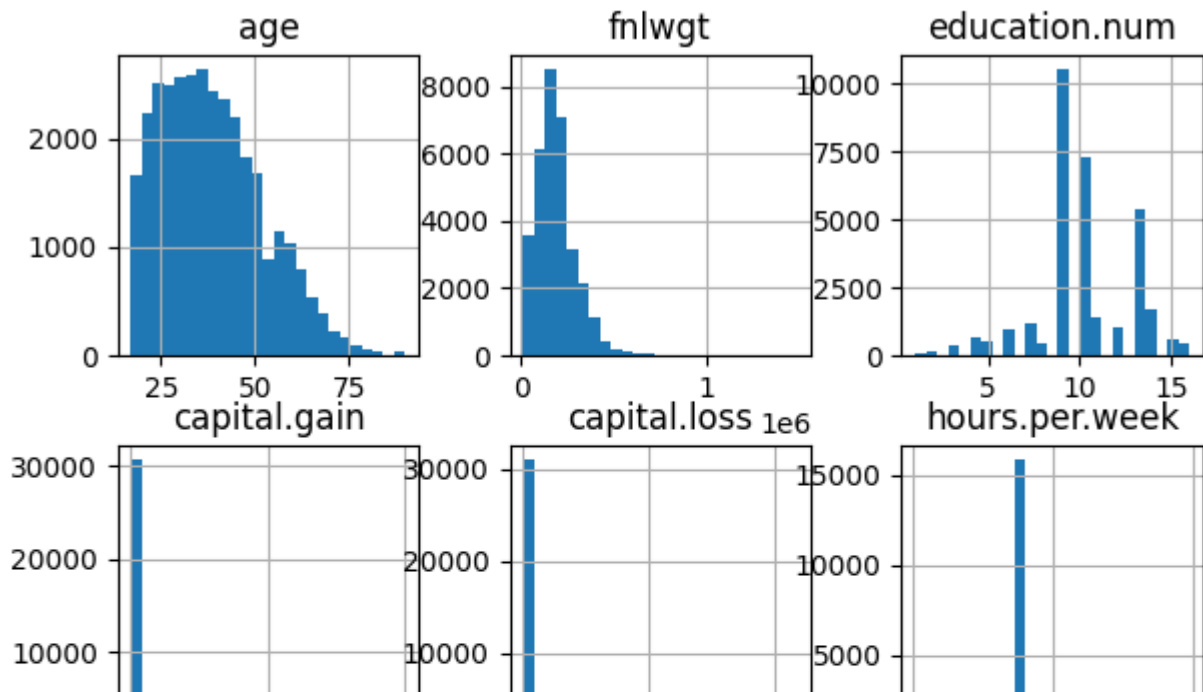
```
df['income'] = df['income'].astype('int')
```

<ipython-input-48-c88e10f6120e>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/user>
df['income'] = df['income'].astype('int')


```
categorical = [col for col in df.columns if df[col].dtype == 'object' ]
numerical = [col for col in df.columns if df[col].dtype != 'object' ]
```

```
df[numerical].hist(bins=25, figsize=(7, 7))
plt.show()
```

```
sns.histplot(df, x='age', hue='income', bins=32)
```

```
<Axes: xlabel='age', ylabel='Count'>
```




```
#describe ages in classes
def age_group(df):
    age_bins = [0, 20, 40, 60, float('inf')]
    age_labels = ['1', '2', '3', '4']
    df_age_range = df.copy()
```

```
df_age_range['age'] = pd.cut(df_age_range['age'], bins=age_bins, labels=age_lab)
return df_age_range
```

```
df = age_group(df).copy()
```

```
sns.histplot(df, x='age', hue='income', bins= 32)
```

```
<Axes: xlabel='age', ylabel='Count'>
```



▼ Missing Values

```
df_missing = (df=='?').sum()
print(df_missing)
```

```
age          0
workclass    1836
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   1843
relationship 0
race         0
sex          0
```

```
capital.gain      0
capital.loss      0
hours.per.week    0
native.country    582
income            0
dtype: int64
```

```
#dropping row having missing values from dataset
```

```
df = df[df['workclass'] != '?']
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']
df.head()
```

| | age | workclass | fnlwgt | education | education.num | marital.status | occupation | relatio |
|---|-----|-----------|--------|--------------|---------------|----------------|-------------------|---------|
| 1 | 4 | Private | 132870 | HS-grad | 9 | Widowed | Exec-managerial | Not-in- |
| 3 | 3 | Private | 140359 | 7th-8th | 4 | Divorced | Machine-op-inspct | Unnr |
| 4 | 3 | Private | 264663 | Some-college | 10 | Separated | Prof-specialty | Ow |
| 5 | 2 | Private | 216864 | HS-grad | 9 | Divorced | Other-service | Unnr |
| 6 | 2 | Private | 150601 | 10th | 6 | Separated | Adm-clerical | Unnr |

```
df_missing = (df=='?').sum()
print(df_missing)
```

```
age            0
workclass      0
fnlwgt         0
education      0
education.num  0
marital.status 0
occupation     0
relationship   0
race           0
sex            0
capital.gain   0
capital.loss   0
hours.per.week 0
native.country 0
income         0
dtype: int64
```

▼ Data Preparation

```
from sklearn import preprocessing
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

| | workclass | education | marital.status | occupation | relationship | race | sex | native |
|---|-----------|--------------|----------------|-------------------|---------------|-------|--------|--------|
| 1 | Private | HS-grad | Widowed | Exec-managerial | Not-in-family | White | Female | Unit |
| 3 | Private | 7th-8th | Divorced | Machine-op-inspct | Unmarried | White | Female | Unit |
| 4 | Private | Some-college | Separated | Prof-specialty | Own-child | White | Female | Unit |
| 5 | Private | HS-grad | Divorced | Other-service | Unmarried | White | Female | Unit |
| 6 | Private | 10th | Separated | Adm-clerical | Unmarried | White | Male | Unit |

```
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

| | workclass | education | marital.status | occupation | relationship | race | sex | native.cou |
|---|-----------|-----------|----------------|------------|--------------|------|-----|------------|
| 1 | 2 | 11 | 6 | 3 | 1 | 4 | 0 | |
| 3 | 2 | 5 | 0 | 6 | 4 | 4 | 0 | |
| 4 | 2 | 15 | 5 | 9 | 3 | 4 | 0 | |
| 5 | 2 | 11 | 0 | 7 | 4 | 4 | 0 | |
| 6 | 2 | 0 | 5 | 0 | 4 | 4 | 1 | |

```
df = df.drop(df_categorical.columns, axis=1)
df = pd.concat([df, df_categorical], axis=1)
df['income'] = df['income'].astype('category')
df.head()
```

| | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | income | workc |
|---|-----|--------|---------------|--------------|--------------|----------------|--------|-------|
| 1 | 4 | 132870 | 9 | 0 | 4356 | 18 | 1 | |
| 3 | 3 | 140359 | 4 | 0 | 3900 | 40 | 1 | |
| 4 | 3 | 264663 | 10 | 0 | 3900 | 40 | 1 | |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30139 entries, 1 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30139 non-null  category
1   fnlwgt                30139 non-null  int64
2   education.num         30139 non-null  int64
3   capital.gain          30139 non-null  int64
4   capital.loss          30139 non-null  int64
5   hours.per.week        30139 non-null  int64
6   income                30139 non-null  category
7   workclass             30139 non-null  int64
8   education             30139 non-null  int64
9   marital.status        30139 non-null  int64
10  occupation            30139 non-null  int64
11  relationship          30139 non-null  int64
12  race                  30139 non-null  int64
13  sex                   30139 non-null  int64
14  native.country        30139 non-null  int64
dtypes: category(2), int64(13)
memory usage: 3.3 MB
```

▼ Splitting dataset

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop('income',axis=1)
X = X.drop('sex',axis=1)
y=df['income']
X.head()
```

| | age | fnlwgt | education.num | capital.gain | capital.loss | hours.per.week | workclass | ed |
|---|-----|--------|---------------|--------------|--------------|----------------|-----------|----|
| 1 | 4 | 132870 | 9 | 0 | 4356 | 18 | 2 | |

```
y.head()
```

```
1    1
3    1
4    1
5    1
6    1
Name: income, dtype: category
Categories (2, int64): [0, 1]
```

```
X_train , X_test , y_train , y_test = train_test_split(X,y,test_size=0.20)
```

▼ Applying Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)
```

▼
DecisionTreeClassifier

DecisionTreeClassifier(max_depth=5)

```
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
```

```
y_pred_default=dt_default.predict(X_test)
print("confusion matrix\n",confusion_matrix(y_test,y_pred_default))
print(classification_report(y_test,y_pred_default))
```

```
confusion matrix
[[ 764  735]
 [ 214 4315]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.51 | 0.62 | 1499 |
| 1 | 0.85 | 0.95 | 0.90 | 4529 |
| accuracy | | | 0.84 | 6028 |
| macro avg | 0.82 | 0.73 | 0.76 | 6028 |
| weighted avg | 0.84 | 0.84 | 0.83 | 6028 |

```
print("accuracy score", accuracy_score(y_test, y_pred_default))
```

accuracy score 0.8425680159256802

✓ 0s completed at 2:45 AM

